

AutoML Agent Project Synopsis

INDEX

S.NO.	TOPIC	PAGE NO.
1.	Introduction	1
2.	Survey of Existing Softwares/Applications/Patents/Copyright	2
3.	Need of the Project	4
4.	Objectives	5
5.	Hardware/Software Requirements (SRS)	6
6.	Methodology (with Data Flow Diagram)	8
7.	Project Planning	12
8.	References	14

1. Introduction

The **AutoML Agent** project aims to democratize Machine Learning (ML) by providing an intuitive, natural language-driven platform. Traditionally, ML development requires deep expertise in data science, programming, and MLOps, limiting widespread adoption. Our agent translates human-readable goals into complex, optimized ML pipelines using cutting-edge AutoML techniques and Large Language Models (LLMs).

Users simply describe their desired outcome in plain English (e.g., "Predict customer churn"). The system then intelligently identifies the ML task, selects the best AutoML engine (Auto-sklearn, PyCaret, H2O.ai), and orchestrates the entire model development lifecycle, including data preprocessing, feature engineering, model selection, hyperparameter tuning, and deployment as an API endpoint.

This project bridges the gap between complex ML methodologies and diverse user needs, empowering individuals and organizations to unlock their data's potential without extensive programming or data science backgrounds. It accelerates

innovation by enabling rapid prototyping and deployment of AI solutions, transforming data into actionable intelligence with unprecedented ease.

2. Survey of Existing Softwares/Applications/Patents/Copyright

The AutoML field has advanced significantly. This section surveys existing commercial platforms and open-source libraries, identifying their features, strengths, and limitations, and positioning the AutoML Agent.

2.1 Commercial AutoML Platforms

These offer end-to-end cloud-based solutions for data preparation, model building, deployment, and monitoring.

- **Google Cloud AutoML:** Specialized for tabular data, vision, NLP. GUI-driven. Scalable, user-friendly, but vendor lock-in and potentially expensive.
- **AWS SageMaker Autopilot:** Automates tabular model building and tuning. Good AWS integration, explainability, but requires AWS knowledge and cost scales.
- **Microsoft Azure Machine Learning (Automated ML):** Visual and code-first options. Strong Azure integration, but can be complex for new users and costly.
- **DataRobot:** Enterprise AI platform with comprehensive AutoML, MLOps, XAI. Feature-rich, high automation, but very expensive.
- **H2O.ai Driverless AI:** Automated feature engineering, validation, deployment. Strong in feature engineering and explainable AI, but proprietary and resource-intensive.

2.2 Open-Source AutoML Libraries

These provide programmatic access for greater flexibility and cost-effectiveness.

- **Auto-sklearn:** Built on scikit-learn, uses Bayesian optimization and ensembling for algorithm and hyperparameter selection. Effective for tabular data, but computationally intensive and resource-heavy.
- **PyCaret:** Low-code library automating ML tasks (data prep, training, deployment) for various ML types. Easy to use, fast prototyping, but offers less fine-grained control.
- **TPOT (Tree-based Pipeline Optimization Tool):** Uses genetic programming to optimize ML pipelines. Explores wide solutions, but can be very slow and computationally expensive.

2.3 LLM Integrations in ML/MLOps

LLMs like Gemini, GPT-series, and Flan-T5 are used for:

- **Code Generation:** Generating ML code snippets or pipelines from natural

language.

- **Data Analysis & Summarization:** Summarizing datasets or ML experiment results.
- **Chatbot Interfaces for MLOps:** Natural language interaction with MLOps systems.

2.4 Competitive Landscape Analysis

Feature/Aspect	Google Cloud AutoML	AWS SageMaker Autopilot	DataRobot	Auto-sklearn	PyCaret	AutoML Agent (Proposed)
Ease of Use (Non-expert)	High (GUI)	High (GUI)	High (GUI)	Low (Code)	Medium (Low-Code)	Very High (NL Interface)
Customizability	Low	Medium	Medium	High	High	High (via advanced options/LLM)
Deployment Automation	Built-in	Built-in	Built-in	Manual	Medium	Fully Automated (API)
Cost	High (Cloud)	High (Cloud)	Very High (SaaS)	Free (Open Source)	Free (Open Source)	Flexible (Cloud compute + API fees)
Core Value	Cloud-native AutoML	Integrated AWS ML	Enterprise AI	Robust automation	Rapid prototyping	NL-driven end-to-end AutoML for all
Key Limitation	Vendor lock-in	AWS ecosystem bias	Expensive/Overkill	Complex setup	Less control	Complexity of integrating layers

Observations	Good for enterprise on GCP.	Good for tabular ML on AWS.	Leader, but costly.	Powerful, for coders.	Fast for coders.	Unique proposition of universal accessibility through NL.
---------------------	-----------------------------	-----------------------------	---------------------	-----------------------	------------------	--

2.5 Patents/Copyright Considerations

The project uses open-source libraries and integrates with commercial LLM APIs.

- **Open-Source Licenses:** Adherence to licenses (MIT, Apache 2.0, BSD) of all incorporated open-source libraries.
- **LLM API Usage:** Compliance with API Terms of Service and Usage Policies (data privacy, rate limits, responsible AI).
- **Proprietary Code:** Unique orchestration logic, natural language parsing, and frontend design will be copyrighted, acknowledging underlying open-source components.

3. Need of the Project

ML adoption faces significant barriers. The AutoML Agent directly addresses these challenges, providing a solution for various stakeholders.

3.1 Problem Context & Market Pain Points

- **Skill Gap:** Shortage of data scientists and ML engineers limits ML adoption.
- **Complexity & Time-Consumption:** The ML lifecycle (data preprocessing, feature engineering, model selection, hyperparameter tuning, deployment) is complex and time-consuming. Data preprocessing alone can take 60-80% of project time.
- **Cost:** Hiring data science teams or using proprietary AutoML platforms is often prohibitive.
- **Lack of Accessibility:** Business users and domain experts lack technical skills to leverage ML, relying on overburdened technical teams.
- **Slow Time-to-Market:** Lengthy ML development cycles hinder rapid business response.

3.2 Quantifying the Pain-Point (Industry Lens)

- **Market Size:** The global AutoML market is projected to reach USD 14.5 billion by 2030 (CAGR of 56.4% from 2023) [1], indicating high demand.

- **Target Segments:** SMEs, Business Analysts, Software Developers, Junior Data Scientists.
- **Revenue Impact:**
 - **Reduced Operational Costs:** 50-80% reduction in time per ML project (e.g., \$4,000 savings per project).
 - **Accelerated Time-to-Insight/Market:** Faster model deployment means quicker iterations and new revenue.
 - **Increased Innovation:** Lowered entry barrier enables more internal ML experimentation.

3.3 Proof of Need

Surveys consistently show "lack of skilled people" and "difficulty integrating ML" as top AI adoption challenges. The rise of "citizen data scientists" validates the market's need for accessible ML. Our project marries advanced AutoML with natural language interaction, offering unique accessibility.

3.4 Value Proposition

The **AutoML Agent** offers:

- **Democratization of ML:** Empowers non-technical users with plain English.
- **Accelerated Development Cycles:** Drastically reduces time from data to deployed model.
- **Cost-Effectiveness:** Minimizes need for expensive ML talent for routine tasks.
- **Optimal Performance:** Leverages leading AutoML engines for best models.
- **Seamless Integration:** Deploys models as readily consumable API endpoints.
- **Focus on Business Value:** Users concentrate on problem definition and results, not technicalities.

4. Objectives

The AutoML Agent project aims to provide an accessible and efficient platform for automated machine learning. We have defined the following SMART objectives:

1. **Objective 1: Develop an Intuitive Natural Language Interface for ML Task Definition.**
 - **Description:** Implement an LLM-powered module for accurately identifying ML task type, target, and optimization metric from user prompts.
 - **KPI:** Achieve $\geq 90\%$ accuracy in identifying task/target from 100 diverse prompts.
 - **Success Metric:** Non-technical users find prompt interface easy and effective.

2. **Objective 2: Enable End-to-End Automated Model Training and Deployment.**
 - **Description:** Integrate and orchestrate multiple AutoML engines (PyCaret, Auto-sklearn) for data preprocessing, feature engineering, model selection, hyperparameter tuning, and seamless RESTful API deployment.
 - **KPI:** Successfully train and deploy API for 100% of tabular datasets ($\leq 50\text{MB}$) within 60 minutes, achieving competitive performance.
 - **Success Metric:** Every completed job provides a functional, accessible API endpoint.
3. **Objective 3: Provide Real-time Status Monitoring and Results Visualization.**
 - **Description:** Implement asynchronous task management and a user-friendly dashboard for real-time progress, key performance metrics, and API endpoint access.
 - **KPI:** Progress bar updates every 3-5 seconds; final metrics/API displayed within 1 minute of completion.
 - **Success Metric:** User satisfaction for job status visibility and result clarity is $\geq 85\%$.
4. **Objective 4: Establish a Secure and Scalable Multi-user Platform.**
 - **Description:** Implement robust Django user authentication/authorization for secure, isolated access to user data. Containerize architecture with Docker/Docker Compose for scalability.
 - **KPI:** Zero critical security vulnerabilities (e.g., SQL injection, XSS) from automated scans. System handles 50 concurrent training jobs with stable performance ($\leq 10\%$ request failure rate).
 - **Success Metric:** User data is securely segregated, and platform performs stably under load.

5. Hardware/Software Requirements (Software Requirements Specification - SRS)

This section outlines essential hardware/software for AutoML Agent development and deployment, including functional and non-functional SRS.

5.1 Software Requirements Specification (SRS)

5.1.1 Functional Requirements (FR)

- **FR-1: User Authentication and Authorization:** Register, login, maintain sessions, logout. Access core features and data strictly for authenticated users (FR-1.1 to FR-1.6).
- **FR-2: Dataset Management:** Upload (CSV, XLSX, XLS, $\leq 50\text{MB}$), securely store, list, and delete user datasets (FR-2.1 to FR-2.6).

- **FR-3: Natural Language Prompt Input:** Provide text input for ML task description. LLM interprets prompt to extract ML parameters (task type, target, metric). (FR-3.1 to FR-3.3).
- **FR-4: Advanced AutoML Configuration (Optional):** Allow users to specify model type, evaluation metric, training time limit, and validation split ratio. User settings override LLM interpretations (FR-4.1 to FR-4.2).
- **FR-5: AutoML Training Initiation:** Initiate async AutoML training using dataset and parsed parameters. System selects engine, performs data prep, feature engineering, model selection, hyperparameter tuning (FR-5.1 to FR-5.4).
- **FR-6: Training Status Monitoring:** Display real-time progress, provide Job ID, notify on completion/failure (FR-6.1 to FR-6.3).
- **FR-7: Model Deployment & Prediction API:** Auto-deploy best model as RESTful API post-training. Provide API URL. API accepts inputs and returns predictions (FR-7.1 to FR-7.3).
- **FR-8: Model Management & Visualization:** List trained models with metadata. (Future scope: detailed metrics/visualizations). Allow model decommissioning/deletion (FR-8.1 to FR-8.3).

5.1.2 Non-Functional Requirements (NFR)

- **NFR-1: Performance:** Frontend load $\leq 3s$. Auth/status API responses $< 200ms$. AutoML training completes within specified time_limit (≤ 60 min for $\leq 50MB$ datasets). Prediction API responses $\leq 100ms$ (NFR-1.1 to NFR-1.4).
- **NFR-2: Security:** Hashed passwords, HTTPS (prod), CSRF/SQLi/XSS prevention. Isolated user data. Secure API key storage (NFR-2.1 to NFR-2.6).
- **NFR-3: Scalability:** Support ≥ 100 concurrent users, ≥ 10 concurrent training jobs. Independent service scaling (NFR-3.1 to NFR-3.3).
- **NFR-4: Reliability:** $\geq 99.5\%$ uptime (prod). Graceful failure handling for AutoML jobs. Durable data storage (NFR-4.1 to NFR-4.3).
- **NFR-5: Usability:** Intuitive UI for non-ML experts. Clear, actionable error messages (NFR-5.1 to NFR-5.2).
- **NFR-6: Maintainability:** Modular, well-commented codebase; requirements.txt for dependencies (NFR-6.1 to NFR-6.2).

5.2 Hardware Requirements

- **Development/Testing:** Quad-core CPU, 16-32GB RAM, 256-\$500+\$GB SSD.
- **Production (Cloud-based):**
 - **Frontend (Django):** 2 vCPU, 4-8GB RAM.
 - **Backend (Flask):** 4-8 vCPU, 16-32GB RAM.
 - **AutoML Workers:** Dedicated instances (4-16 vCPU, 32-64GB RAM, optional

GPU).

- **Database (PostgreSQL):** Managed service, appropriately sized.
- **Object Storage:** S3/GCS.

5.3 Software Requirements (Specific to Tech Stack)

- **Operating System:** Linux (Ubuntu/Debian), macOS, Windows.
- **Containerization:** Docker Engine, Docker Compose.
- **Programming Languages:** Python 3.9+
- **Frontend Framework:** Django 5.0+ (djangorestframework).
- **Backend Framework:** Flask 2.3+ (Flask-CORS).
- **AutoML Libraries:** auto-sklearn, pycaret, h2o (client).
- **LLM Integration:** openai, google-generativeai, transformers, torch/tensorflow.
- **Database:** PostgreSQL 14+ (prod), SQLite 3 (dev).
- **Task Queue:** Celery 5.x (Redis broker, PostgreSQL result backend).
- **WSGI HTTP Server:** Gunicorn 21+.
- **Web Server (Proxy):** Nginx (prod).

5.4 Bill of Materials / Licenses / Compliance Checklist

- **Bill of Materials (Conceptual):** Cloud Compute, Managed DB, Object Storage, LLM API Credits, Domain, SSL.
- **Licenses:** Adherence to open-source licenses (MIT, Apache 2.0, BSD), Django (BSD-3-Clause), Flask (BSD 3-Clause), and commercial LLM API terms.
- **Compliance Checklist (Initial):** Data Privacy (GDPR/CCPA), Security (audits, access control, OWASP Top 10), Responsible AI (bias mitigation, explainability), Monitoring.

6. Methodology

The AutoML Agent project will follow an **Agile development methodology** (Scrum), emphasizing iterative development, continuous feedback, adaptability, and frequent delivery.

6.1 Agile Methodology Overview

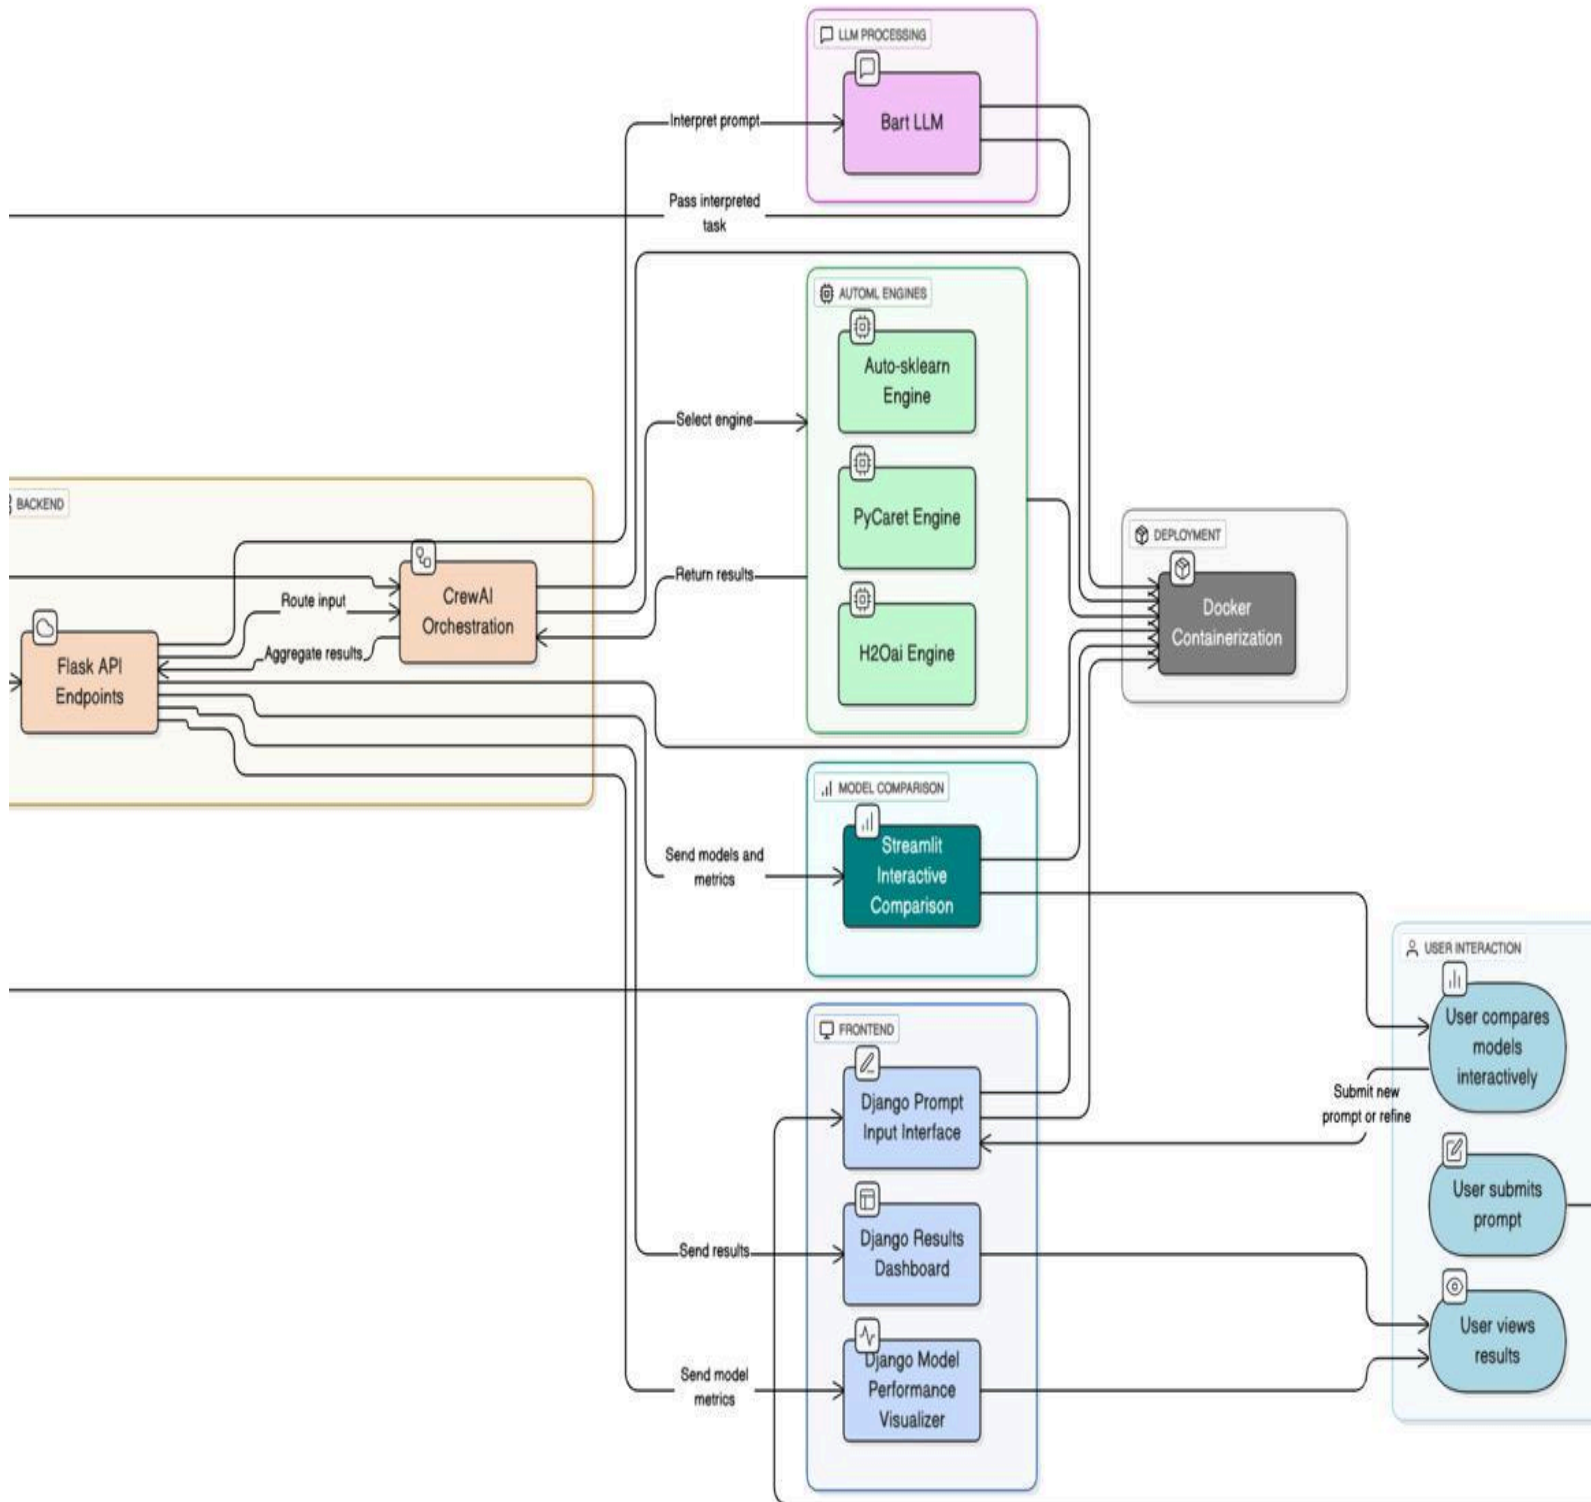
- **Sprints:** 2-week iterations for vertical slices of functionality.
- **Backlog:** Prioritized product backlog.
- **Daily Scrums:** Short daily syncs for progress and impediments.
- **Sprint Review:** End-of-sprint demonstration to stakeholders.
- **Sprint Retrospective:** Team reflection for process improvement.
- **CI/CD:** Foster early integration and frequent testing.

6.2 Solution Architecture

Microservices-oriented architecture: distinct frontend, backend, and asynchronous worker layer, orchestrated by Docker.

- **Frontend (Django):** Web interface, user sessions, sends requests to Flask.
- **Backend (Flask API Service):** Receives requests, LLM interaction, dispatches AutoML jobs, provides status, serves deployed models.
- **LLM Layer:** (Within Flask) Integrates external LLM APIs (OpenAI/Gemini) for prompt interpretation.
- **AutoML Engines:** (Within Celery workers) Libraries (Auto-sklearn, PyCaret, H2O.ai) for core ML logic.
- **Asynchronous Task Queue (Celery):** Manages long-running AutoML jobs (Redis broker, PostgreSQL results).
- **Database (PostgreSQL):** Stores user accounts, dataset metadata, job statuses, model details.
- **Object Storage (e.g., AWS S3):** Scalable storage for raw datasets and model artifacts.
- **Deployment (Docker, Gunicorn, Docker Compose):** Containerizes services, Gunicorn serves applications, Docker Compose orchestrates.

6.3 Data Flow Diagram (DFD)



Data Flow Explanation:

1. **User Interaction:** User accesses Django Frontend.
2. **Authentication:** Frontend sends credentials to Django Auth API, which interacts with Django Auth System.
3. **Dataset Upload:** User uploads dataset via Frontend to Flask Dataset API, saved to Object Storage.
4. **Prompt & Training:** User inputs prompt to Frontend, sent to Flask ML Train API.
5. **LLM Parsing:** ML Train API sends prompt to LLM Parser, extracting AutoML parameters.
6. **AutoML Orchestration:** Flask's AutoML Orchestrator receives params and dataset ref.
7. **Job Dispatch:** Orchestrator dispatches job to Celery Broker (Redis).
8. **Async Processing:** Celery Workers pick up jobs.
9. **AutoML Execution:** Worker loads dataset, runs AutoML Engine, saves Model Artifact to Object Storage.
10. **Status Update & Model Info:** Worker updates job status in PostgreSQL. Model deployment event triggers Model Serving Service.
11. **Model Deployment:** Model Serving Service loads artifact, exposes API, stores API URL in PostgreSQL.
12. **Status Polling:** Frontend polls job status from PostgreSQL (via Flask API).
13. **Display Results:** Frontend displays status and API endpoint.
14. **Model Consumption:** User consumes deployed Model API for predictions.

6.4 Development Workflow

- **Sprint Planning:** Define tasks from product backlog.
- **Coding & Unit Testing:** Write code with comprehensive unit tests.
- **Containerization:** Dockerize new services early.
- **Integration Testing:** Test services with docker-compose.
- **Version Control:** Git with clear branching.
- **Code Reviews:** Peer review all code changes.
- **Documentation:** Incremental technical and user docs.

- **Feedback Loop:** Continuous stakeholder feedback.

7. Project Planning

This outlines the roadmap, milestones, and release plan, including risk management.

7.1 Roadmap & Release Plan

Phase 1: Core MVP (Week 1-2)

- **Goal:** Foundational architecture, secure auth, basic NL prompt, simulated AutoML, real-time status.
- **Milestones:** Functional Django Frontend (M1.1), Standalone Flask Backend (M1.2), Frontend-Backend comms (M1.3), Basic LLM parsing (M1.4).
- **Deliverables:** Django web app, Flask API, Docker setup, user auth, simulated ML.

Phase 2: Feature Expansion & Core AutoML Integration (week 3-4)

- **Goal:** Real AutoML engine, async task queue, persistent storage, model deployment.
- **Milestones:** Celery integration (M2.1), PyCaret integration (M2.2), Dataset/model persistence (M2.3), Model serving API (M2.4).
- **Deliverables:** Real AutoML training, background jobs, persistent storage, deployable model APIs.

Phase 3: Beta & Refinement (Week 5-6)

- **Goal:** Enhance AutoML, improve UI/UX, robust error handling, public beta prep.
- **Milestones:** Secondary AutoML integration (M3.1), Advanced error handling (M3.2), UI/UX improvements (M3.3), Internal Beta (M3.4).
- **Deliverables:** Enhanced AutoML, robust error reporting, polished UI.

Phase 4: General Availability (GA) & Continuous Improvement (week 7+)

- **Goal:** Public launch, monitoring, ongoing feature development.
- **Milestones:** Public Beta (M4.1), Monitoring/Alerting (M4.2), Full GA Launch (M4.3).
- **Deliverables:** Live, stable, publicly accessible AutoML Agent.

7.2 Budget & Resource Allocation (Conceptual)

- **Personnel:** Project mentor, Jovac team members (team of 4)
- **Compute & Storage:** Cloud credits (VMs, DB, object storage).
- **API Costs:** LLM API usage fees.
- **Tools & Licenses:** Dev tools, commercial licenses.

7.3 Risk Register & Mitigation

Risk	Description	Likelihood	Impact	Mitigation Strategy
Data Quality Issues	Messy, incomplete, or inappropriate user data.	High	High	Robust data validation/preprocessing. Clear data format guidelines. Graceful error handling.
LLM Interpretation Accuracy	LLM misinterprets prompts.	Medium	Medium	Clear prompt engineering. User confirmation/override. Strong foundational LLM.
AutoML Engine Performance/Compatibility	Engines perform poorly or have integration issues.	Medium	High	Benchmark multiple engines. Develop easy-swap interfaces. Test with diverse datasets.
Scalability & Concurrency	System overwhelmed by concurrent training jobs.	Medium	High	Async task queues (Celery). Containerization /orchestration (Docker Compose, Kubernetes). Resource monitoring.
Security Vulnerabilities	Exposed APIs, insecure data storage.	High	High	Secure coding (CSRF, input validation, auth). Regular audits. Least privilege.
Cost Overruns	Unexpected	Medium	Medium	Cost

	cloud/API costs.			monitoring/alerts. Resource optimization. Reserved instances.
Talent Availability	Difficulty finding expertise.	Medium	Medium	Cross-training. Open-source community leverage. Clear documentation.

8. References

This lists academic papers, industry whitepapers, and relevant standards.

[1] Grand View Research. (2024). *AutoML Market Size, Share & Trends Analysis Report By Component (Platform, Services), By Deployment (On-Premise, Cloud), By End-use, By Region, And Segment Forecasts, 2024 - 2030*. [Online]. Available: <https://www.grandviewresearch.com/industry-analysis/automl-market> (Accessed: June 26, 2025).

[2] Feurer, M., Eggensperger, K., & Hutter, F. (2020). Auto-sklearn: Efficient and Robust Automated Machine Learning. *Journal of Machine Learning Research*, 21(44), 1-33. [Online]. Available: <https://www.jmlr.org/papers/v21/19-089.html>

[3] Ali, M., et al. (2020). PyCaret: An Open-Source, Low-Code Machine Learning Library in Python. *arXiv preprint arXiv:2012.15570*. [Online]. Available: <https://arxiv.org/abs/2012.15570>

[4] Lytle, M., et al. (2022). H2O Driverless AI. In: *Data Science Best Practices*. Springer, Cham.

[5] M. S. Hassan, et al. (2023). A Survey on Large Language Models for Software Engineering. *arXiv preprint arXiv:2303.04870*. [Online]. Available: <https://arxiv.org/abs/2303.04870>

[6] ISO/IEC 27001:2022. (2022). *Information security, cybersecurity and privacy protection - Information security management systems - Requirements*. International Organization for Standardization.

[7] O'Reilly Media. (2020). *The AI Maturity Model: A Guide for Data Leaders*. [Online].

Available: <https://www.oreilly.com/content/the-ai-maturity-model/> (Accessed: June 26, 2025).